ISA5810 DM2026 Lab 1 — Master Notebook Questions

Fill in this file as you work through DM2026-Lab1-Master.ipynb. Submit this completed file converted to PDF together with the rest of your Lab 1 work.

Please do not change the style of this document.

Question numbers 1–10 follow the same order as the notebook (top to bottom).

Student Information

Please complete the fields below. For **English Name** and **Chinese Name**, you may fill in **only one** if you do not have the other.

- **English Name:**
- **Chinese Name:蕭博丞**
- **Email:damian50503@gmail.com**
- **Student ID:112204048**
- **University:National ChengChi University**

Instructions

These questions go with the Master Jupyter notebook (DM2026-Lab1-Master.ipynb). Each question is tied to a specific notebook section. Open the notebook, find **Question N (take home)**, and write your answer here.

We **strongly** suggest running the notebook to learn the lab, but the TAs will **not** audit whether you executed every cell to answer the questions in this document.

How to answer each question

1. Find the **notebook section** listed under the question (use the Table of Contents if needed).
2. Read the task; change parameters when asked (e.g. stop_words, thresholds, k, minSup).
3. Write a clear answer in the **Answer** space below each question.
4. Leave the question text unchanged; write only in the answer and disclosure fields.

Sources, honesty, and AI use

Acceptable vs. Unacceptable AI Usage: Generative AI should serve as an interactive tutor to guide your understanding, not as an automated solver completing the assignment for you.

- **Permitted (AI as Guide & Editor):**
  - Asking AI to explain unfamiliar theoretical concepts, discuss algorithmic trade-offs, or troubleshoot runtime errors in the notebook.
  - Polishing the grammar, phrasing, and readability of your own original, self-written drafts.

- **Prohibited (AI as Solver):**
  - Copying assignment questions directly into an AI tool to generate answers.
  - Using AI to fabricate experimental results or observations without executing the code in DM2026-Lab1-Master.ipynb.
  - Submitting machine-generated prose wholesale.

- **External material:** If you use websites, books, papers, or documentation beyond the notebook, cite them with **clickable links** (or DOI/ISBN when no URL exists).

- **When AI was used:** Your answer is acceptable if you (1) **disclose** AI use and (2) include **at least one external link** that validates the information in your answer (official docs, textbook, paper, or reputable tutorial — not the AI chat alone as the only source).

- **When AI was not used:** You may write "No" under **AI used** or leave it blank only if you truly did not use AI for that question. References are still encouraged if you looked anything up.

Under each question, complete:

- **AI used for this answer:** Yes / No
- **If yes, how:** (e.g. brainstorming, grammar, concept explanation)
- **External references (links):** *(required when AI = Yes; encouraged otherwise)*

Write in English. Keep answers concise and specific.

Questions

Question 1: Data Filtering & Domain Relevance

Notebook: Section 3.2 — Question 1 (take home)

Task: Filter the dataframe **X** for the category **sci.med** and check its length.

Questions:

1. How many documents are in this specific category?
2. From a data mining perspective, why is isolating specific categories a crucial first step before applying algorithms like frequent pattern mining?

Answer: 1.0 2.Filtering a specific category removes unrelated data, so frequent pattern mining can find clearer and more relevant patterns.

AI used for this answer: No

If yes, how:

External references (links):

Question 2: Handling Missing Data

Notebook: Section 4.1 — Question 2 (take home) (after the **dropna()** demonstration)

Task: In Section 4.1, a dummy record with a missing value is injected and then removed using the **dropna()** function.

Question: If 40% of documents were missing the **category_name** label, why would simply dropping these rows (with **dropna()**) be a flawed data mining strategy? Name one objective alternative so the text data can still be utilized.

Answer: Simply using dropna() would remove 40% of the documents, so we would lose a lot of useful text and the remaining data might become biased. One alternative is to train a classifier using the labeled documents and use it to predict the missing labels. This allows us to keep and use the text data.

AI used for this answer: No

If yes, how:

External references (links):

Question 3: Data Sampling Risks

Notebook: Section 5.1 — Question 3 (take home)

Task: In Section 5.1, the code uses **X_sample = X.sample(n=1000)** to create a subset. Compare the bar chart of the original category distribution to the sampled distribution.

Questions:

1. Do the exact category proportions perfectly match the original?
2. What is the specific risk of simple random sampling on a highly imbalanced dataset?

Answer: 1.No, the category proportions are similar, but they do not perfectly match because the sample is selected randomly.

2.In a highly imbalanced dataset, simple random sampling may select too few or none of the minority-category records. This can make the sample biased and cause the model to perform poorly on minority categories.

AI used for this answer: No

If yes, how:

External references (links):

Question 4: Text Vectorization and Matrix Sparsity

Notebook: Section 5.3 — Question 4 (take home)

Task: In Section 5.3, the **CountVectorizer** converts text into a document-term matrix using **stop_words='english'**. Change this parameter to **stop_words=None** and re-run the cell to check the shape of **X_counts**.

Questions:

1. Does the total number of features (columns) increase or decrease?
2. Since stopwords are highly common words that appear in almost all documents, what effect does adding them back have on the overall sparsity (the percentage of zeros) of the matrix?

Answer: 1.After changing stop_words='english' to stop_words=None, the number of features increased from **35,482 to 35,788** because stopwords were added back.

2.The sparsity decreased because common stopwords appear in many documents, so the matrix contains more non-zero values and a lower percentage of zeros.

AI used for this answer: No

If yes, how:

External references (links):

Question 5: Feature Selection — Variance Thresholding

Notebook: Section 5.3.1 — Question 5 (take home)

Task: In the variance filtering section, re-run **variance_filter** using threshold values **threshold = 0.0, threshold = 0.01, and threshold = 0.05**. Record the number of features marked as **keep == True**.

Questions:

1. What happens to the total size of the vocabulary as the variance threshold increases from 0.0 to 0.05?
2. In terms of document frequency, what kinds of words have variance close to zero and are eliminated first?
3. Why is eliminating near-zero variance features considered a safe way to reduce dimensionality without harming model accuracy?

Answer:

Threshold	**Number of terms kept (keep == True)**
0.0	35,482
0.01	5,089
0.05	1,491

AI used for this answer: No

If yes, how:

External references (links):

Question 6: Feature Correlation — Linear (Pearson) vs. Monotonic (Spearman)

Notebook: Section 5.3.2 — Question 6 (take home)

Task: Run the correlation filter against the binary indicator for the **comp.graphics** class. Inspect the top 10 features sorted by **pearson_r** and compare them to the top 10 features sorted by **spearman_r**.

Questions:

1. Are the top-correlated words generic stop words or domain-specific terminology for graphics?
2. What is the mathematical difference in what **Pearson** correlation measures versus what **Spearman** rank correlation measures?
3. In what specific scenario would a word have a high **Spearman** rank correlation with a topic, but a significantly lower **Pearson** correlation?

Answer: 1.Most of the words are related to computer graphics, such as graphics, image, 3d, video, and vga. There are also some general words like thanks and looking, but common stop words were already removed.

2.Pearson measures the linear relationship using the original word counts. Spearman uses the ranking of the counts, so it can also find relationships that are not perfectly linear.

3.A word may have high Spearman but lower Pearson when it usually appears more often in graphics documents, but its exact counts are uneven. For example, most graphics documents may use the word a few times, while one document uses it many times. This affects Pearson more than Spearman.

AI used for this answer: No

If yes, how:

External references (links):

Question 7: Frequent Pattern Mining — Algorithm & Parameter Dynamics

Notebook: Section 5.4.2 — Question 7 (take home)

Task: In the Frequent Pattern Mining section, run **Top-K FAE** with **k = 100 and k = 500**. Then, run **MaxFPGrowth** using minimum support thresholds **minSup=9 and minSup= 3**.

Questions:

1. How does decreasing **minSup** (from 9 to 3) in **MaxFPGrowth** affect the execution runtime and the total number of discovered patterns?
2. How does the parameter k in **Top-K FAE** differ fundamentally from **minSup** in **MaxFPGrowth** when deciding which patterns to keep?
3. Why are "maximal" frequent patterns structurally more concise than standard frequent patterns?

Answer (1) — Runtime and pattern count when lowering minSup:

When minSup decreased from 9 to 3, the number of patterns increased from 710 to 9,547. The runtime also increased because the algorithm had to search through many more patterns.

Answer (2) — Top-K k vs MaxFPGrowth minSup:

In Top-K FAE, k directly decides how many patterns to keep. For example, k=100 keeps 100 patterns. In MaxFPGrowth, minSup sets the minimum number of times a pattern must appear, so the final number of patterns is not fixed.

Answer (3) — Maximal vs standard frequent patterns:

Maximal patterns are more concise because smaller patterns are not included if they are already part of a larger frequent pattern. This reduces repeated and similar results.

AI used for this answer: No

If yes, how:

External references (links):

Question 8: Dimensionality Reduction — Motivation

Notebook: Section 5.5 — Question 8 (take home) (before PCA/t-SNE code)

Task: After text vectorization, the **X_counts** matrix creates tens of thousands of distinct features (words).

Questions:

1. If we want to visually map how these documents group by category on a scatter plot, why is it mathematically and visually impossible to plot the raw **X_counts** matrix directly?
2. What specific category of data mining algorithms (referenced in the helper files) must be applied first?

Answer: 1.The raw X_counts matrix has more than 35,000 features because each word is one dimension. A normal scatter plot can only display two or three dimensions, so we cannot plot the raw matrix directly.

2.We need to use dimensionality reduction algorithms first. For example, PCA, t-SNE, or UMAP can reduce the data to two dimensions so it can be shown on a scatter plot.

AI used for this answer: No

If yes, how:

External references (links):

Question 9: Dimensionality Reduction — Linear (PCA) vs. Non-linear Manifolds (t-SNE/UMAP)

Notebook: Section 5.5 — Question 9 (take home) (after 2D/3D plots)

Task: In the 2D/3D Dimensionality Reduction section, compare the scatter plot projection generated by **PCA** against that generated by **t-SNE** or **UMAP**.

Questions:

1. Which method **(PCA or t-SNE/UMAP)** shows distinct, separated cluster islands for the 4 document categories, and which shows heavy overlap around the center?
2. Why does **PCA** struggle to separate high-dimensional text categories compared to **t-SNE** or **UMAP**?
3. Why must data analysts be careful not to over-interpret the physical distance between separate clusters in a **t-SNE** plot?

Answer:1.t-SNE and UMAP show clearer groups than PCA, although some categories still overlap. PCA has heavy overlap near the center and only a few outliers.

2.PCA only finds linear directions with the largest variance. Text data is sparse and may contain non-linear relationships, so PCA cannot separate the categories well. t-SNE and UMAP are better at keeping nearby documents together.

3.We should not over-interpret the distance between clusters in t-SNE because it mainly preserves local relationships. Two clusters looking far apart does not always mean they are very different in the original data.

AI used for this answer: Yes / No

If yes, how:

External references (links):

Question 10: One-Hot Binarization

Notebook: Section 5.6 — Question 10 (take home)

Task: In Section 3.1, examine **X['category']** (integers 0, 1, 2, 3). Then look at the output of **LabelBinarizer** in Section 5.6, where each category is converted into a binary vector (e.g., [1, 0, 0, 0]).

Questions:

1. How many binary columns are created to represent the 4 categories?
2. If an algorithm (like linear regression or KNN) is given the raw integer column **(0, 1, 2, 3)** as input, what false mathematical assumption does it make about the categories?
3. How does one-hot binarization eliminate this false assumption?

Answer:1.

Four binary columns are created because there are four categories.
2.Using the numbers 0, 1, 2, 3 directly makes the algorithm think the categories have an order and meaningful distance. For example, it may assume category 3 is greater than category 1, but the category numbers are only labels.
3.One-hot encoding gives each category its own column. Only the correct category has a value of 1, so there is no false ranking between the categories.
AI used for this answer: No

If yes, how:

External references (links):
